

# Introduction to Data Science, Topic 5

- Instructor: Professor Henry Horng-Shing Lu,  
Institute of Statistics, National Chiao Tung University, Taiwan  
Email: [hslu@stat.nctu.edu.tw](mailto:hslu@stat.nctu.edu.tw)
- WWW: <http://www.stat.nctu.edu.tw/misg/hslu/course/DataScience.htm>
- Reference:  
M. A. Pathak, Beginning Data Science with R, 2014, Springer-Verlag.
- Evaluation: Homework: 50%, Term Project: 50%
- Office hours: By appointment

# Course Outline

- Introduction of Data Science
- Introduction of R
- Process Real Data by R
- Data Visualization
- **Exploratory Data Analysis**
- Regression
- Classification
- Text Mining
- Clustering

# Exploratory Data Analysis with R

References:

Ch. 5, M. A. Pathak, Beginning Data Science with R, 2014, Springer-Verlag.





An approximate answer to the right  
problem is worth a good deal more  
than an exact answer to an  
approximate problem.

— *John Tukey* —

AZ QUOTES

[http://www.azquotes.com/author/14847-John\\_Tukey](http://www.azquotes.com/author/14847-John_Tukey)

John W. Tukey wrote the book *Exploratory Data Analysis* in 1977. Tukey held that too much emphasis in statistics was placed on statistical hypothesis testing (confirmatory data analysis); more emphasis needed to be placed on using data to suggest hypotheses to test.

# Exploratory Data Analysis

- ***Summary Statistics***
- ***Getting Sense of Data Distribution***
- ***Putting it all together: Outlier Detection***

In this topic, we'll use the US Census Demographic Data on Kaggle:

[https://www.kaggle.com/muonneutrino/us-census-demographic-data/data#\\_=\\_](https://www.kaggle.com/muonneutrino/us-census-demographic-data/data#_=_)

# Summary Statistics – Data size

the `dim()` function output the numbers of rows and columns.

```
> census <- read.csv("acs2015_county_data.csv")
> dim(census)
[1] 3220    37
```

Our data contain 3220 records (in this census data means 3220 counties), and 37 variables.

We can also use `nrow()` and `ncol()` function to find only the number of rows and columns, respectively.

```
> nrow(census)
[1] 3220
> ncol(census)
[1] 37
```

# Summary Statistics – Summarizing the Data

The `head()` and `tail()` functions, output the first and last few entries of a data frame. It is useful to have a glimpse of data.

```
> head(census[, 1:8])
```

	CensusId	State	County	TotalPop	Men	Women	Hispanic	White
1	1001	Alabama	Autauga	55221	26745	28476	2.6	75.8
2	1003	Alabama	Baldwin	195121	95314	99807	4.5	83.1
3	1005	Alabama	Barbour	26932	14497	12435	4.6	46.2
4	1007	Alabama	Bibb	22604	12073	10531	2.2	74.5
5	1009	Alabama	Blount	57710	28512	29198	8.6	87.9
6	1011	Alabama	Bullock	10678	5660	5018	4.4	22.2

# Summary Statistics – Summarizing the Data

The `summary()` functions give the brief summary for each column.

```
> summary(census[, 1:8])
```

CensusId	State	County	TotalPop
Min. : 1001	Texas : 254	Washington: 31	Min. : 85
1st Qu.:19033	Georgia : 159	Jefferson : 26	1st Qu.: 11218
Median :30024	Virginia: 133	Franklin : 25	Median : 26035
Mean :31394	Kentucky: 120	Jackson : 24	Mean : 99409
3rd Qu.:46106	Missouri: 115	Lincoln : 24	3rd Qu.: 66430
Max. :72153	Kansas : 105	Madison : 20	Max. :10038388
	(Other) :2334	(Other) :3070	

Men	Women	Hispanic	White
Min. : 42	Min. : 43	Min. : 0.000	Min. : 0.00
1st Qu.: 5637	1st Qu.: 5572	1st Qu.: 1.900	1st Qu.:64.10
Median : 12932	Median : 13057	Median : 3.900	Median :84.10
Mean : 48897	Mean : 50512	Mean :11.012	Mean :75.43
3rd Qu.: 32993	3rd Qu.: 33488	3rd Qu.: 9.825	3rd Qu.:93.20
Max. :4945351	Max. :5093037	Max. :99.900	Max. :99.80



# Summary Statistics – Summarizing the Data

For categorical variables like State and County, the summary contains the number of times occur in each value, so we can see that Washington occur 31 times. We can see all Washington records by:

```
> census[which(census$County == "Washington"), 1:6]
  CensusId      State      County TotalPop    Men   Women
65      1129      Alabama Washington   16997    8490    8507
183      5143      Arkansas Washington  216432  108144  108288
306      8121      Colorado Washington    4795    2482    2313
387     12133      Florida Washington   24629   13478   11151
537     13303      Georgia Washington   20785   10467   10318
595     16087       Idaho Washington   10025    5082    4943
690     17189      Illinois Washington   14457    7242    7215
785     18175      Indiana Washington   27930   13867   14063
881     19183       Iowa  Washington   22017   10847   11170
.
.
.
```

# Summary Statistics – Summarizing the Data

For numeric variables, the summary contain:

- Min.—smallest value of the variable.
- 1st Qu. (Q1)—first quartile or 25th percentile
- Median—second quartile or 50th percentile
- Mean—Average value of the variable.
- 3rd Qu. (Q2)—third quartile or 75th percentile
- Max.—largest value of the variable.

These statistics are useful to get a sense of the data distribution for a variable: its range and centrality.

# Summary Statistics – Ordering Data by a Variable

`sort ( )` function sorts vectors or data frame by a variable.

```
> sort(census$TotalPop)
 [1]      85      117      267      433      443      448      548      551      565      565
[11]     606     643     673     675     681     705     711     733     756     769
[21]     776     778     781     812     820     821     847     851     874     901
...

```

`sort ( )` function sorts vectors in descending order, if we set `decreasing = T`

```
> sort(census$TotalPop, decreasing = T)
 [1] 10038388  5236393  4356362  4018143  3223096  3116069
 [7]  2639042  2595259  2485003  2301139  2298032  2094769
[13]  2045756  2035572  1914526  1868149  1843152  1825502
...

```

`sort ( )` function sorts also can sort string alphabetically.

# Summary Statistics – Ordering Data by a Variable

`order( )` function orders the data frame for given variable in one step.

```
> census[order(-census$TotalPop)[1:3], 1:6]
```

	CensusId	State	County	TotalPop	Men	Women
205	6037	California	Los Angeles	10038388	4945351	5093037
611	17031	Illinois	Cook	5236393	2537245	2699148
2624	48201	Texas	Harris	4356362	2166727	2189635

The minus sign in front of the `census$TotalPop` means sorted in reverse. `order( )` function can also sort on multiple variables, like `order(variable1, variable2, ...)`.

# Summary Statistics – Group and Split Data by a Variable

We can select a subset by:

```
> which(census$State == "California")
```

But if we want to do this repeatedly and perform the analysis on that, we can use `by( )` function.

```
> by(census$TotalPop, census$State, mean)
census$State: Alabama
[1] 72098.81
-----
census$State: Alaska
[1] 25288.79
-----
...
```

# Summary Statistics – Group and Split Data by a Variable

There is a similar function `split()`, which split the data and output it in a list.

```
> data.split = split(census, census$State)
> for (x in names(data.split)) {
+   dd = data.split[[x]]
+   print(x)
+   print(dd[order(-dd$TotalPop)[1:5], c(3, 4)])
+ }
[1] "Alabama"
      County TotalPop
37  Jefferson  659026
49   Mobile   414251
45   Madison  346438
51 Montgomery 228138
59   Shelby  203530
...
```

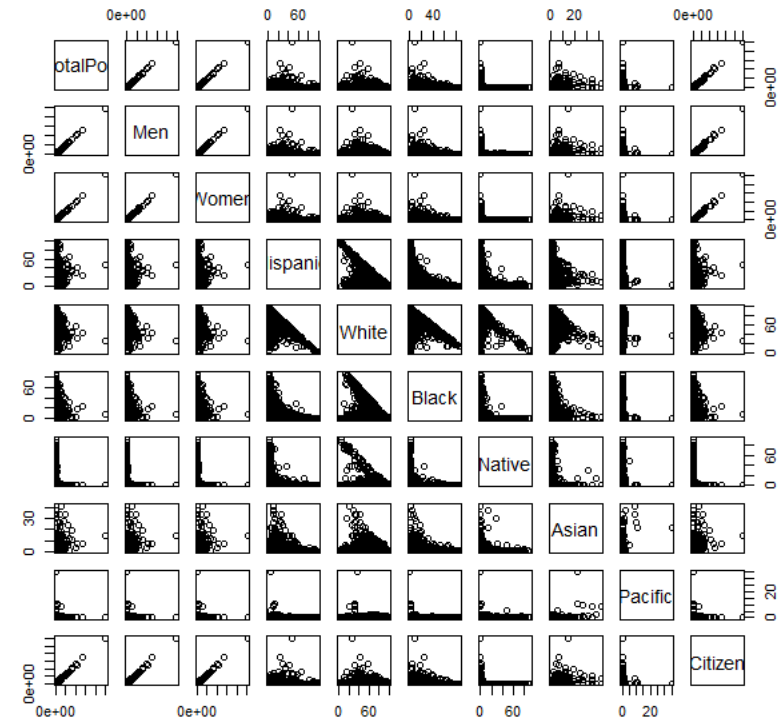
# Summary Statistics – Variable Correlation

`cor()` function computes the correlation of a pair of variable.

```
> cor(census$TotalPop, census$Men)
[1] 0.9998772
```

We can use the `cor()` function to obtain the pairwise correlation between a set of numeric variables

```
> cor(census[, 4:13])
> pairs(census[, 4:13])
```



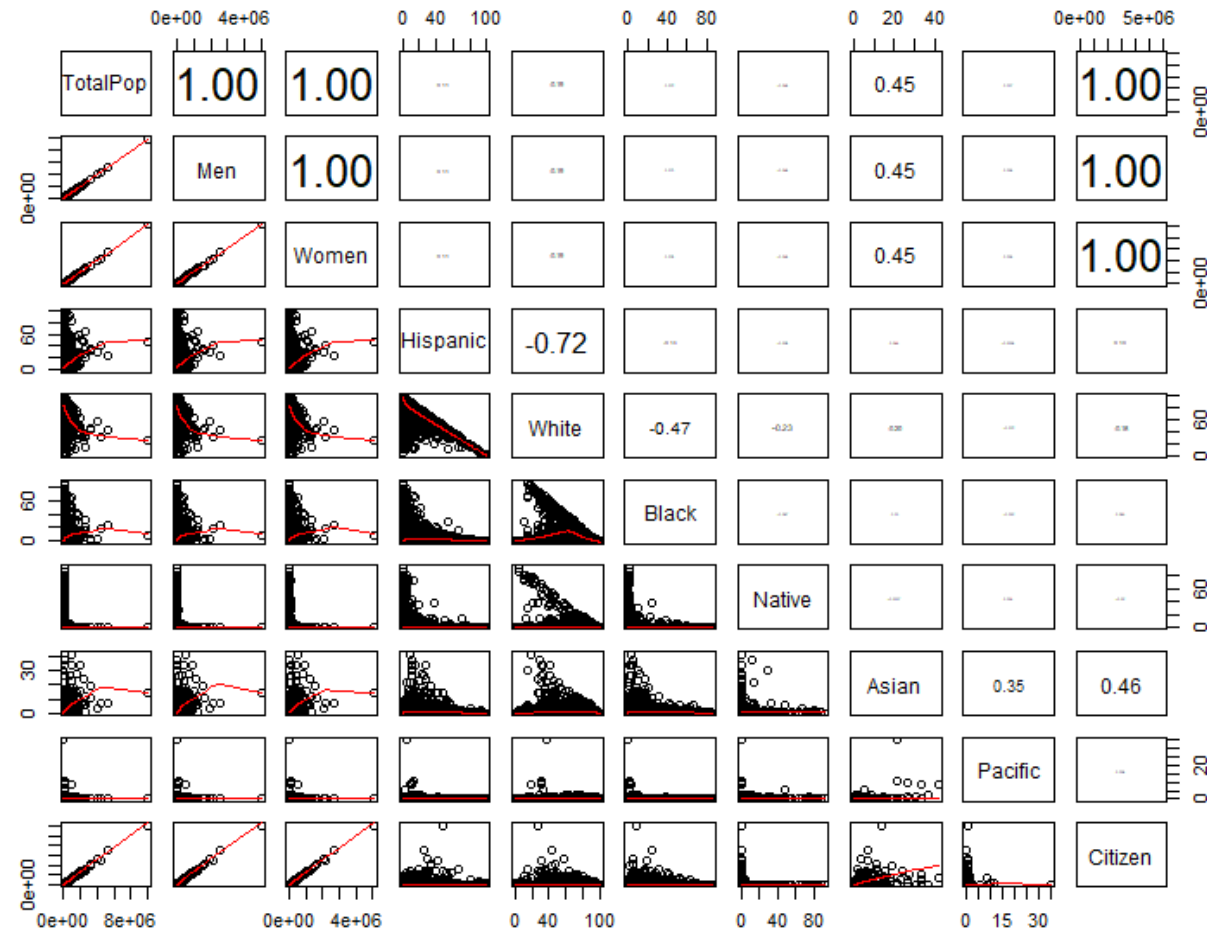
# Summary Statistics – Variable Correlation

Or we can customized pairs plot:

```
> panel.cor <- function(x, y, digits = 2, prefix = "", cex.cor =  
10, ...) {  
+   usr <- par("usr"); on.exit(par(usr))  
+   par(usr = c(0, 1, 0, 1))  
+   r <- cor(x, y)  
+   txt <- format(c(r, 0.123456789), digits = digits)[1]  
+   txt <- paste(prefix, txt, sep = "")  
+   if(missing(cex.cor)) cex.cor <- 1/strwidth(txt)  
+   text(0.5, 0.5, txt, cex = cex.cor * abs(r))  
+ }  
> pairs(census[, 4:13], lower.panel = panel.smooth,  
+       upper.panel = panel.cor)
```



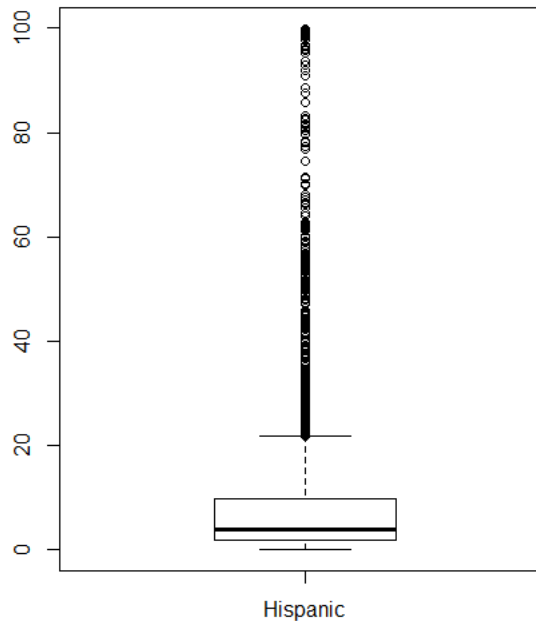
# Summary Statistics – Variable Correlation



# Getting a Sense of Data Distribution— Box Plots

Box plots based on the five-number summary statistics of a variable (minimum, Q1, median, Q3, maximum)

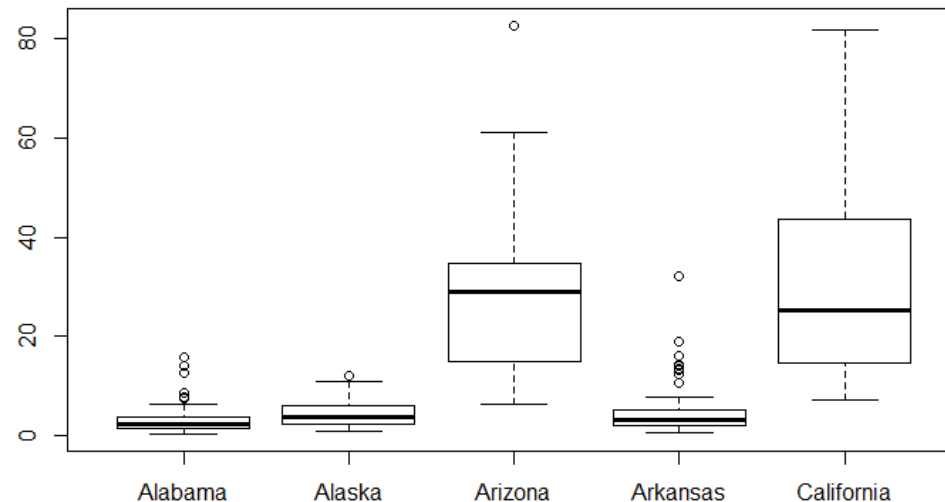
```
boxplot(census$Hispanic, names = c("Hispanic"), show.names = T)
```



# Getting a Sense of Data Distribution– Box Plots

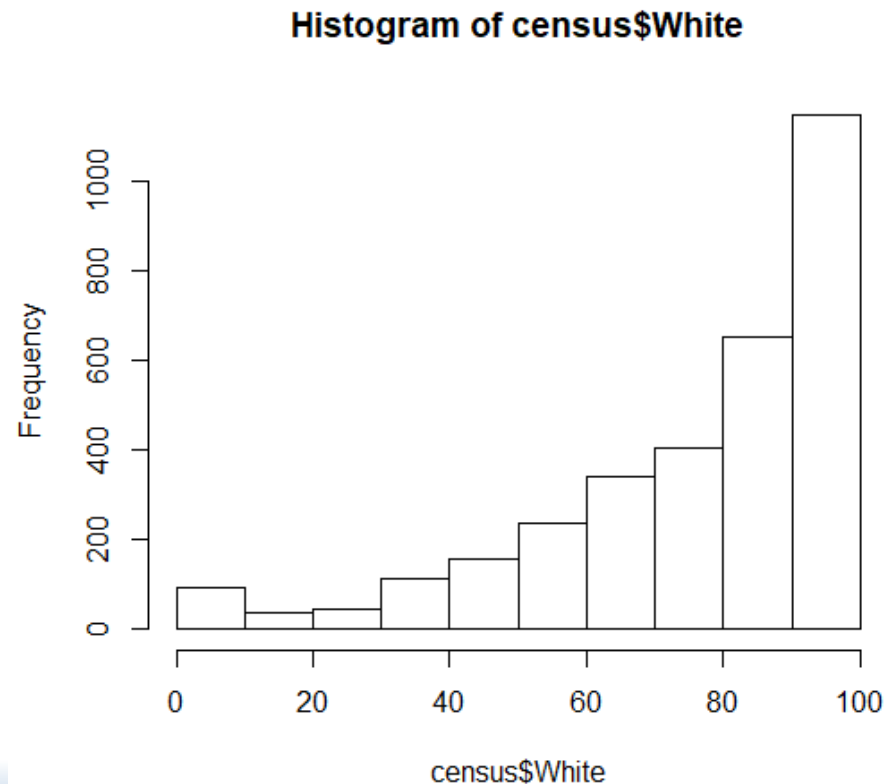
We also can draw boxplot against different States.

```
> census_micro <- subset(census, census$State %in%  
unique(census$State)[1:5])  
> census_micro$State <- as.factor(as.character(census_micro$State))  
> boxplot(census_micro$Hispanic ~ census_micro$State)
```



# Getting a Sense of Data Distribution— Histograms

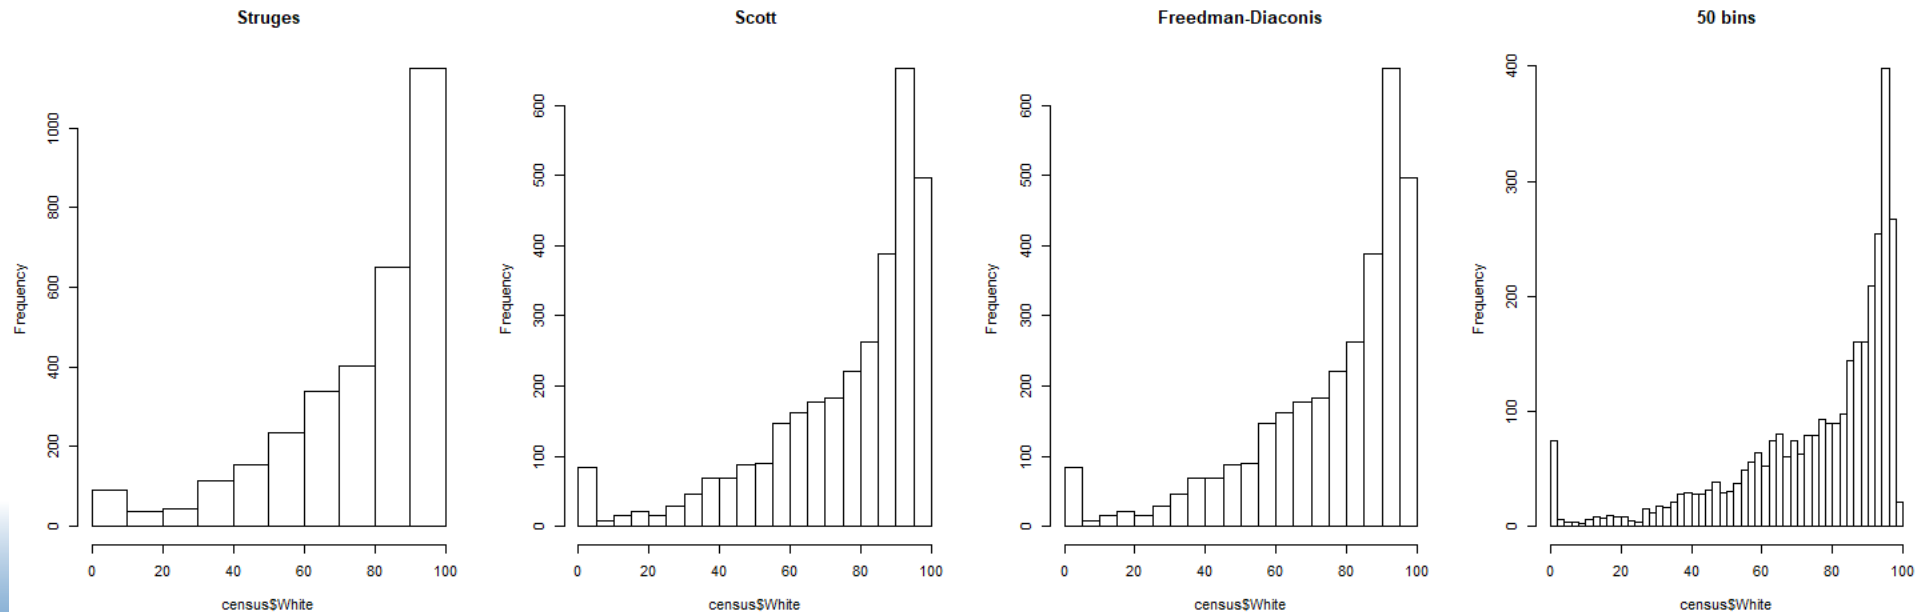
```
> hist(census$White)
```



# Getting a Sense of Data Distribution—Histograms

We also can use different break.

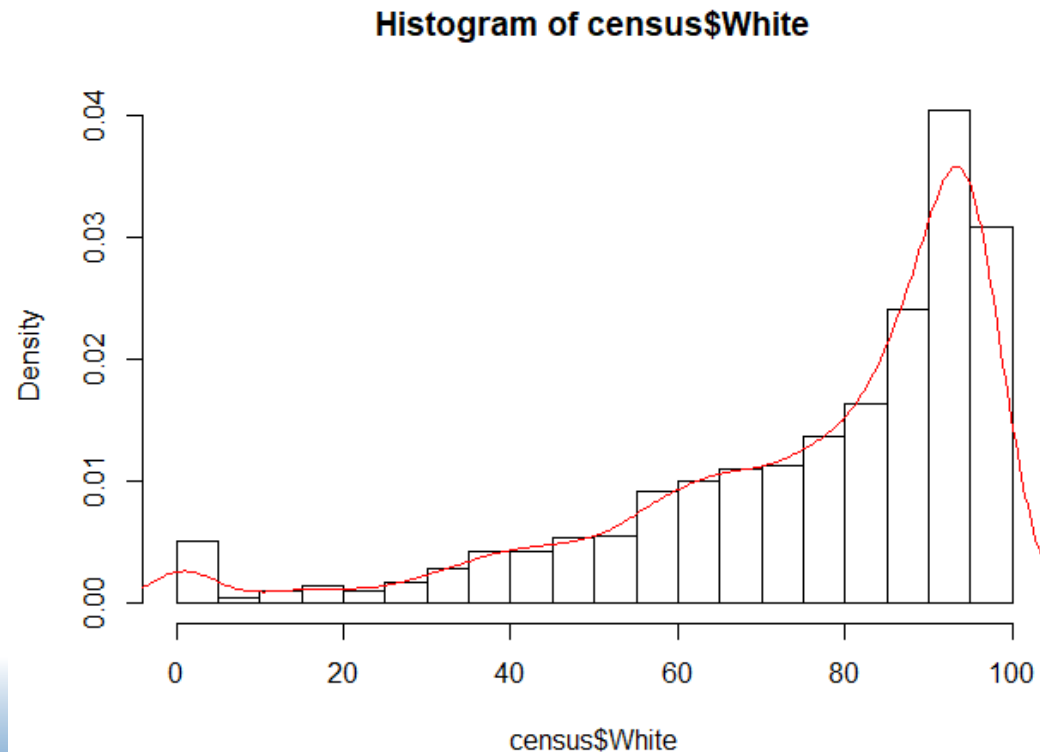
```
> par(mfrow = c(1, 4))  
> hist(census$White, breaks = "sturges", main = "Struges")  
> hist(census$White, breaks = "scott", main = "Scott")  
> hist(census$White, breaks = "fd", main = "Freedman-Diaconis")  
> hist(census$White, breaks = 50, main = "50 bins")
```



# Getting a Sense of Data Distribution— Histograms

Add density to histogram.

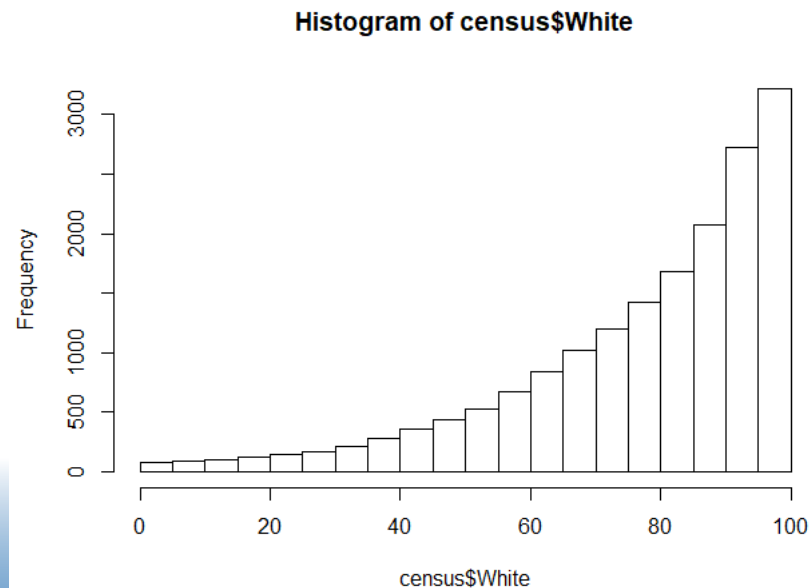
```
> hist(census$White, breaks = "FD", freq = F)  
> points(density(census$White), type = "l", col = "red")
```



# Getting a Sense of Data Distribution– Histograms

Cumulative Histogram.

```
> cumhist = function(x) {  
+   h = hist(census$White, "FD", plot = F)  
+   h$counts = cumsum(h$counts)  
+   plot(h)  
+ }  
> cumhist(census$White)
```



# Getting a Sense of Data Distribution– Measuring Data Symmetry Using Skewness and Kurtosis

Skewness, Kurtosis and Gini.

```
> library(moments)
> skewness(census[, 4:13])
  TotalPop      Men      Women  Hispanic      White      Black
14.287924 14.395679 14.182423   3.216816 -1.431692   2.321071
  Native      Asian  Pacific  Citizen
 8.053653   7.169789 37.178879 12.652094
> kurtosis(census[, 4:13])
  TotalPop      Men      Women  Hispanic      White
344.368589 348.934737 339.835909 13.610745   4.667022
  Black      Native      Asian  Pacific  Citizen
 8.368248   76.003223  75.057308 1676.754633 277.026915
> library(reldist)
> gini(census$TotalPop)
[1] 0.7510696
```



# Getting a Sense of Data Distribution— Measuring Data Symmetry Using Skewness and Kurtosis

Skewness, Kurtosis and Gini.

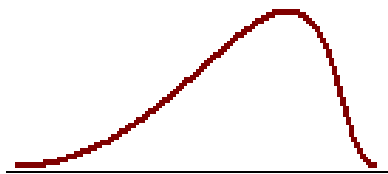
Skewness is a measure of symmetry. (<https://en.wikipedia.org/wiki/Skewness>)

Kurtosis is a measure of whether the data are heavy-tailed or light-tailed relative to a normal distribution. (<https://en.wikipedia.org/wiki/Kurtosis>)

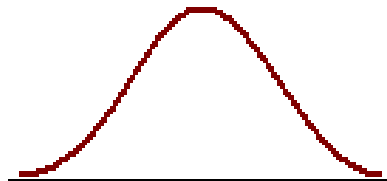
For univariate data  $Y_1, Y_2, \dots, Y_N$ , the formula for skewness and kurtosis are:

$$skewness = \frac{\sum_{i=1}^N (Y_i - \bar{Y})^3 / N}{s^3}, \quad kurtosis = \frac{\sum_{i=1}^N (Y_i - \bar{Y})^4 / N}{s^4}$$

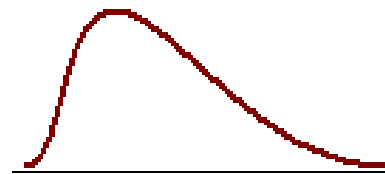
where  $\bar{Y}$  is the mean,  $s$  is the standard deviation, and  $N$  is the number of data points.



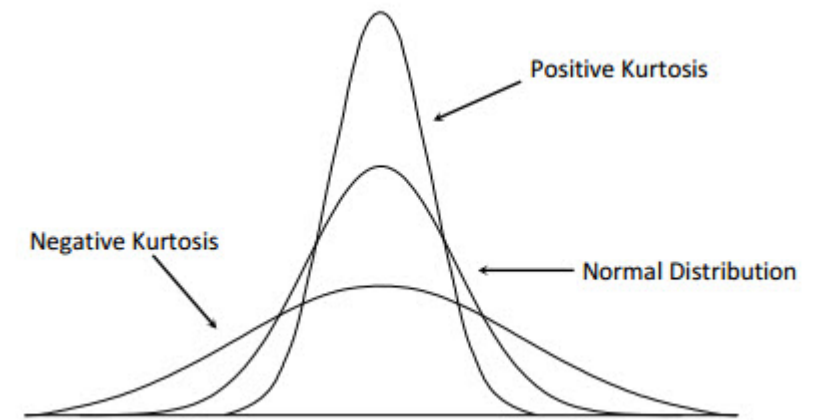
Negatively skewed distribution  
or Skewed to the left  
Skewness < 0



Normal distribution  
Symmetrical  
Skewness = 0



Positively skewed distribution  
or Skewed to the right  
Skewness > 0



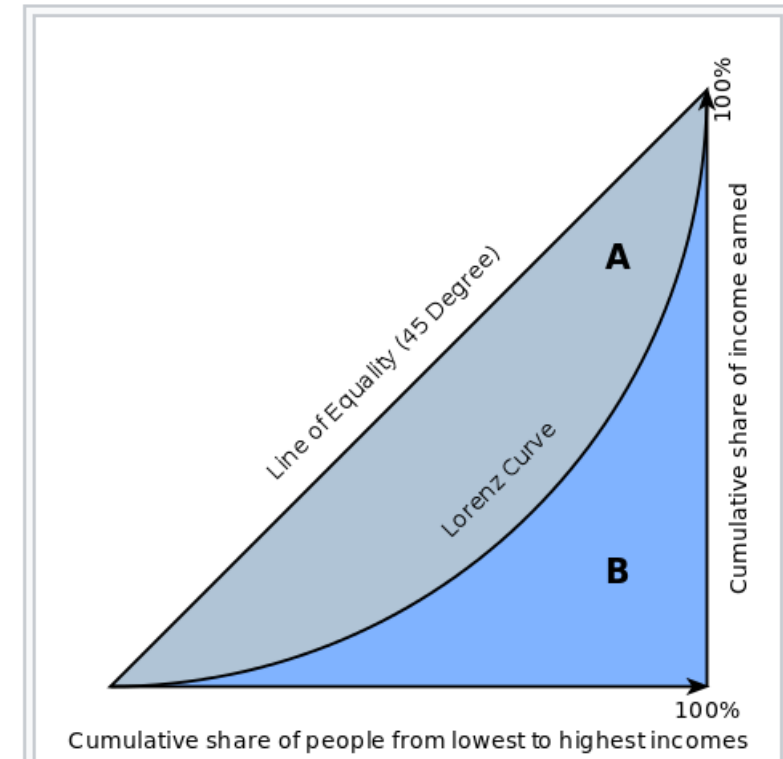
# Getting a Sense of Data Distribution— Measuring Data Symmetry Using Skewness and Kurtosis

Gini coefficient:

The Gini coefficient measures the inequality among values of a frequency. A Gini coefficient of zero expresses perfect equality, where all values are the same. A Gini coefficient of 1 expresses maximal inequality among values.

([https://en.wikipedia.org/wiki/Gini\\_coefficient](https://en.wikipedia.org/wiki/Gini_coefficient))

$$Gini = \frac{\sum_{i=1}^n \sum_{j=1}^n |x_i - x_j|}{2n \sum_{i=1}^n x_i}$$



Graphical representation of the Gini coefficient

The graph shows that the Gini coefficient is equal to the area marked *A* divided by the sum of the areas marked *A* and *B*, that is,  $Gini = A / (A + B)$ . It is also equal to  $2A$  and to  $1 - 2B$  due to the fact that  $A + B = 0.5$  (since the axes scale from 0 to 1).

# Putting It All Together: Outlier Detection

```
> health <- read.csv("ehresp_2014.csv")
> summary(health[, 22:29])
```

eusnap		eugenhth		eugroshp		euhgt	
Min.	:-3.000	Min.	:-3.000	Min.	:-3.000	Min.	:-3.00
1st Qu.:	2.000	1st Qu.:	2.000	1st Qu.:	1.000	1st Qu.:	63.00
Median :	2.000	Median :	2.000	Median :	1.000	Median :	66.00
Mean :	1.868	Mean :	2.477	Mean :	1.503	Mean :	65.63
3rd Qu.:	2.000	3rd Qu.:	3.000	3rd Qu.:	2.000	3rd Qu.:	70.00
Max.	: 2.000	Max.	: 5.000	Max.	: 3.000	Max.	:77.00

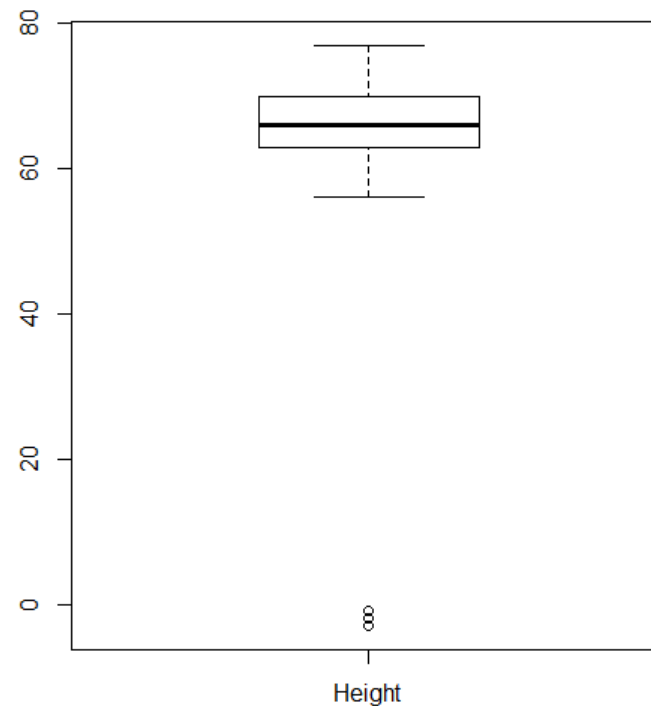
euinclvl		euincome2		eumeat		eumilk	
Min.	:5.000	Min.	:-3.0000	Min.	:-2.0000	Min.	:-3.000
1st Qu.:	5.000	1st Qu.:	-1.0000	1st Qu.:	-1.0000	1st Qu.:	-1.000
Median :	5.000	Median :	-1.0000	Median :	1.0000	Median :	2.000
Mean :	5.177	Mean :	-0.2313	Mean :	0.5293	Mean :	1.158
3rd Qu.:	5.000	3rd Qu.:	1.0000	3rd Qu.:	1.0000	3rd Qu.:	2.000
Max.	:6.000	Max.	: 3.0000	Max.	: 2.0000	Max.	: 2.000

It seems that there is something abnormal in “euhgt”, the min value is -3, but height can’t be negative.

# Putting It All Together: Outlier Detection

Beside using `summary()` to observe the data, we also can detect if there is an outlier by drawing some graph.

```
> boxplot(health$euhtgt, names = c("Height"), show.names = T)
```



# *Interactive Visualizations Using Shiny*

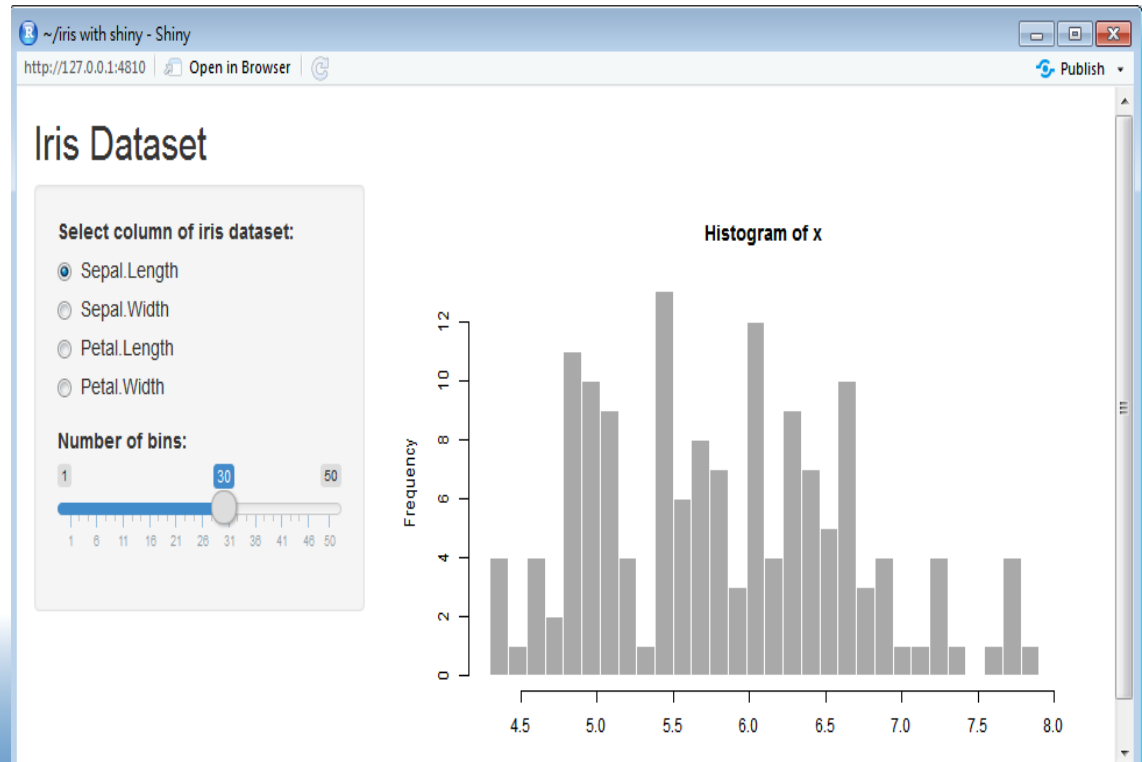
- *UI*
  - library(shiny)
  - shinyUI(fluidPage(
    - #fluid page for dynamically adapting to screens of different resolutions.
    - titlePanel("Iris Dataset"),
    - sidebarLayout(
      - sidebarPanel(
        - #implementing radio buttons
        - radioButtons("p", "Select column of iris dataset:",
          - list("Sepal.Length"='a', "Sepal.Width"='b', "Petal.Length"='c', "Petal.Width"='d')),
        - #slider input for bins of histogram
        - sliderInput("bins",
          - "Number of bins:",
          - min = 1,
          - max = 50,
          - value = 30)
      - # Show a plot of the generated distribution
    - ),
    - mainPanel(
      - plotOutput("distPlot")
    - )
    - )
    - ))

# *Interactive Visualizations Using Shiny*

- **Server**
  - library(shiny)
  - #writing server function
  - shinyServer(function(input, output) {
  - #referring output distPlot in ui.r as output\$distPlot
  - output\$distPlot <- renderPlot({
  - #referring input p in ui.r as input\$p
  - if(input\$p=='a'){
  - i<-1
  - }
  - if(input\$p=='b'){
  - i<-2
  - }
  - if(input\$p=='c'){
  - i<-3
  - }
  - if(input\$p=='d'){
  - i<-4
  - }
  - x <- iris[, i]
  - #referring input bins in ui.r as input\$bins
  - bins <- seq(min(x), max(x), length.out = input\$bins + 1)
  - #producing histogram as output
  - hist(x, breaks = bins, col = 'darkgray', border = 'white')
  - })
  - }

# *Interactive Visualizations Using Shiny*

- Save to R file UI.R and Server.R into same folder
- Execute the command
  - `runApp("folder path")`
  - Ex: `runApp("C:/Users/USER/Documents/iris with shiny/")`



# Demo video

- <https://www.youtube.com/watch?v=bVY804VA5ak>



# *3D dynamic plots with iris*

- `install.packages(c("rgl", "car"))`
- `library(car)`
- `attach(iris)`
- `scatter3d(x = iris$Sepal.Length,`
- `y = iris$Sepal.Width,`
- `z = iris$Petal.Length)`
- `scatter3d(x = iris$Sepal.Length,`
- `y = iris$Sepal.Width,`
- `z = iris$Petal.Length,`
- `groups = iris$Species)`
- `scatter3d(x = iris$Sepal.Length,`
- `y = iris$Sepal.Width,`
- `z = iris$Petal.Length,`
- `groups = iris$Species,`
- `surface=FALSE, ellipsoid = TRUE)`

# Demo video

- <https://www.youtube.com/watch?v=6oFg0tulAxU>
- <https://www.youtube.com/watch?v=ZQGjJFvDSXY>
- <https://www.youtube.com/watch?v=glgkaFAJoGE>

# Homework

- Basic
  - Find a dataset you want to analysis.
  - Do EDA on this dataset, like summary statistics, box plot and histogram...
  - Detect if there have any outlier in this dataset.
- Advanced
  - If there have any outlier in this dataset, how would you deal with it and why?
  - Give your point of view what you found in this dataset.

# Homework 5 (submitted to e3.nctu.edu.tw before Oct 15, 2019)

- Use R and/or other software to visualize the data set with missing data (NA) that you select
- Explain the results you obtain
- Discuss possible problems you plan to investigate for future studies
- Possible source of open data:

UCI Machine Learning Repository

(<http://archive.ics.uci.edu/ml/datasets.html>)

# *References*

1. [https://en.wikipedia.org/wiki/Exploratory\\_data\\_analysis#Development](https://en.wikipedia.org/wiki/Exploratory_data_analysis#Development)
2. [https://www.kaggle.com/muonneutrino/us-census-demographic-data/data#\\_=\\_](https://www.kaggle.com/muonneutrino/us-census-demographic-data/data#_=_)
3. <https://www.kaggle.com/bls/eating-health-module-dataset/data>
4. [https://en.wikipedia.org/wiki/Box\\_plot](https://en.wikipedia.org/wiki/Box_plot)
5. Christie, M. (2001). The Ozone layer: A philosophy of science perspective. United Kingdom: Cambridge University Press. (Chap. 6).
6. Dorfman, R. (1979). A formula for the Gini coefficient. The review of economics and statistics. The Review of Economics and Statistics, 61, 146–149.
7. Most major U.S. cities show population declines. USA Today, June 2011.
8. Size 8 is the new 7: Why our feet are getting bigger. Time Magazine, Oct 2012.
9. Sugary sodas high in diabetes-linked compound.  
<http://abcnews.go.com/Health/Healthday/story?id=4508420&page=1#.UUzdKFt34eF>. March 2007.
10. To his credit, charge card king doesn't cash in. Los Angeles Times, Dec 2004.